

SC3000: Artificial Intelligence

Lab 1 Report: TDDb Team Armaan & Friends

Members:

Goel Armaan (U2123642H)

Implementation of task2, task3 and task3_alt

Chiam Da Jie (U2121233G)

Implementation of all helper functions, task1_UCS and task1_A*

Guo Yuan Lin (U2122626F)

Researching, testing and implementing heuristics, weight_tuning and task3_weighted_energy

We assert that everyone has contributed equally to the creation of this report.

Date of Submission: October 11, 2023

References

- [1] A. Patel, “Heuristics” *Amit’s A* Pages*, 2020. [Online]. Available: <http://theory.stanford.edu/~amitp/GameProgramming/Heuristics.html>. [Accessed: 03-Oct-2023]
- [2] S. Sengupta and B. An, “Exercise 7: Maze Runner.” NTU SC1015, Singapore, Mar-2021.

1 Heuristics

Since we have the $\{x, y\}$ coordinates for each node, we can apply various heuristic functions to estimate the distance between a node and the goal to make use of informed search algorithms such as A*. The following coordinate-based functions will be used as heuristics for this assignment.

1.1 Euclidean Distance

Euclidean distance, also known as straight line distance, measures the length of the line that connects any 2 points on the $\{x, y\}$ plane. The formula is

$$\sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$$

1.2 Manhattan Distance

Manhattan distance assumes movement is possible in 4 directions (up, down, left, right) and calculates the number of units any 2 points are away on a standard square grid. The formula is

$$|x_2 - x_1| + |y_2 - y_1|$$

1.3 Chebyshev Distance

Chebyshev Distance assumes diagonal movement is possible and takes the value of the maximum difference over any of the axis values. Instead of adding the 2 differences like in Manhattan, we take only the maximum value here. The formula is

$$\max(|x_2 - x_1|, |y_2 - y_1|)$$

1.4 Octile Distance

Octile Distance is similar to Chebyshev as it also assumes diagonal movement. But instead of just taking the maximum, it also considers the minimum difference with a weight of $\sqrt{2}$. The formula is

$$\max(|x_2 - x_1|, |y_2 - y_1|) + \sqrt{2} \cdot \min(|x_2 - x_1|, |y_2 - y_1|)$$

1.5 Weight Tuning

From hereon, ‘weight tuning’ refers to the function implemented in *utils.py*. This function takes in the A* task to be executed along with the optimal result obtained from the UCS implementation of the same. It multiplies the heuristic value by weights ranging from 1 to 0, in steps of 0.01 for all 4 heuristics ($w * h(n)$). It crosschecks whether the distance returned is the same as the UCS result and then returns the weight with the least number of node expansions for each heuristic. The results and discussions of running weight tuning are mentioned in their respective task implementation sections.

2 Tasks

2.1 Task 1

Result:

Shortest path: $S \rightarrow 1363 \rightarrow 1358 \rightarrow 1357 \rightarrow 1356 \rightarrow 1276 \rightarrow 1273 \rightarrow 1277 \rightarrow 1269 \rightarrow 1267 \rightarrow 1268 \rightarrow 1284 \rightarrow 1283 \rightarrow 1282 \rightarrow 1255 \rightarrow 1253 \rightarrow 1260 \rightarrow 1259 \rightarrow 1249 \rightarrow 1246 \rightarrow 963 \rightarrow 964 \rightarrow 962 \rightarrow 1002 \rightarrow 952 \rightarrow 1000 \rightarrow 998 \rightarrow 994 \rightarrow 995 \rightarrow 996 \rightarrow 987 \rightarrow 986 \rightarrow 979 \rightarrow 980 \rightarrow 969 \rightarrow 977 \rightarrow 989 \rightarrow 990 \rightarrow 991 \rightarrow 2369 \rightarrow 2366 \rightarrow 2340 \rightarrow 2338 \rightarrow 2339 \rightarrow 2333 \rightarrow 2334 \rightarrow 2329 \rightarrow 2029 \rightarrow 2027 \rightarrow 2019 \rightarrow 2022 \rightarrow 2000 \rightarrow 1996 \rightarrow 1997 \rightarrow 1993 \rightarrow 1992 \rightarrow 1989 \rightarrow 1984 \rightarrow 2001 \rightarrow 1900 \rightarrow 1875 \rightarrow 1874 \rightarrow 1965 \rightarrow 1963 \rightarrow 1964 \rightarrow 1923 \rightarrow 1944 \rightarrow 1945 \rightarrow 1938 \rightarrow 1937 \rightarrow 1939 \rightarrow 1935 \rightarrow 1931 \rightarrow 1934 \rightarrow 1673 \rightarrow 1675 \rightarrow 1674 \rightarrow 1837 \rightarrow 1671 \rightarrow 1828 \rightarrow 1825 \rightarrow 1817 \rightarrow 1815 \rightarrow 1634 \rightarrow 1814 \rightarrow 1813 \rightarrow 1632 \rightarrow 1631 \rightarrow 1742 \rightarrow 1741 \rightarrow 1740 \rightarrow 1739 \rightarrow 1591 \rightarrow 1689 \rightarrow 1585 \rightarrow 1584 \rightarrow 1688 \rightarrow 1579 \rightarrow 1679 \rightarrow 1677 \rightarrow 104 \rightarrow 5680 \rightarrow 5418 \rightarrow 5431 \rightarrow 5425 \rightarrow 5424 \rightarrow 5422 \rightarrow 5413 \rightarrow 5412 \rightarrow 5411 \rightarrow 66 \rightarrow 5392 \rightarrow 5391 \rightarrow 5388 \rightarrow 5291 \rightarrow 5278 \rightarrow 5289 \rightarrow 5290 \rightarrow 5283 \rightarrow 5284 \rightarrow 5280 \rightarrow T$
 Shortest distance: 148648.63722140007

Nodes Expanded:

UCS – 5304; A* – 608 using Octile distance with a weight of 1.00

Weight Tuning:

Heuristic	Weight	Expanded	Reduction (%)
Euclidean	1.00	1225	76.9
Manhattan	0.87	626	88.2
Chebyshev	1.00	1958	63.1
Octile	1.00	608	88.5

Discussion:

The relaxed version of the NYC problem is a shortest path problem, which can be solved with Uniform-Cost Search as we have links of varying non-negative costs, and UCS guarantees optimality under this condition. We have used Python's *heapq* to implement the priority queue. The cost function is set to the distance between 2 nodes. We maintain a dictionary of path distances and a set of explored nodes. In each iteration, a node is popped from the queue. If the popped node already exists in the explored set, it is ignored. If the popped node is the goal state, then we stop. Otherwise, all the neighbours of the popped node are added to the priority queue if they do not already exist in the distance dictionary, or if the new distance is less than the existing one. To increase performance, we also implemented an A* version of Task 1 that uses Octile distance as the heuristic as it results in the least number of nodes expanded. The only difference between the UCS and A* versions is that when adding a node to the queue in A*, the priority is the sum of the total path and heuristic distance. ($g(n) + h(n)$). The NYC problem involves movements in all directions, in which case Euclidean should have been the best. But in real life, we do not have straight paths from every node to the other, so a deconstruction of axes is needed, which all the other heuristic functions do. Octile accounts for diagonal movement, and unlike Chebyshev, also assigns a

weight to the axis with the smaller difference, which could indicate why it has performed the best for this scenario. Manhattan also came close, but at a weight less than 1, most likely because it only assumes 4-directional movement.

2.2 Task 2

Result:

Shortest path: $S \rightarrow 1363 \rightarrow 1358 \rightarrow 1357 \rightarrow 1356 \rightarrow 1276 \rightarrow 1273 \rightarrow 1277 \rightarrow 1269 \rightarrow 1267 \rightarrow 1268 \rightarrow 1284 \rightarrow 1283 \rightarrow 1282 \rightarrow 1255 \rightarrow 1253 \rightarrow 1260 \rightarrow 1259 \rightarrow 1249 \rightarrow 1246 \rightarrow 963 \rightarrow 964 \rightarrow 962 \rightarrow 1002 \rightarrow 952 \rightarrow 1000 \rightarrow 998 \rightarrow 994 \rightarrow 995 \rightarrow 996 \rightarrow 987 \rightarrow 988 \rightarrow 979 \rightarrow 980 \rightarrow 969 \rightarrow 977 \rightarrow 989 \rightarrow 990 \rightarrow 991 \rightarrow 2465 \rightarrow 2466 \rightarrow 2384 \rightarrow 2382 \rightarrow 2385 \rightarrow 2379 \rightarrow 2380 \rightarrow 2445 \rightarrow 2444 \rightarrow 2405 \rightarrow 2406 \rightarrow 2398 \rightarrow 2395 \rightarrow 2397 \rightarrow 2142 \rightarrow 2141 \rightarrow 2125 \rightarrow 2126 \rightarrow 2082 \rightarrow 2080 \rightarrow 2071 \rightarrow 1979 \rightarrow 1975 \rightarrow 1967 \rightarrow 1966 \rightarrow 1974 \rightarrow 1973 \rightarrow 1971 \rightarrow 1970 \rightarrow 1948 \rightarrow 1937 \rightarrow 1939 \rightarrow 1935 \rightarrow 1931 \rightarrow 1934 \rightarrow 1673 \rightarrow 1675 \rightarrow 1674 \rightarrow 1837 \rightarrow 1671 \rightarrow 1828 \rightarrow 1825 \rightarrow 1817 \rightarrow 1815 \rightarrow 1634 \rightarrow 1814 \rightarrow 1813 \rightarrow 1632 \rightarrow 1631 \rightarrow 1742 \rightarrow 1741 \rightarrow 1740 \rightarrow 1739 \rightarrow 1591 \rightarrow 1689 \rightarrow 1585 \rightarrow 1584 \rightarrow 1688 \rightarrow 1579 \rightarrow 1679 \rightarrow 1677 \rightarrow 104 \rightarrow 5680 \rightarrow 5418 \rightarrow 5431 \rightarrow 5425 \rightarrow 5424 \rightarrow 5422 \rightarrow 5413 \rightarrow 5412 \rightarrow 5411 \rightarrow 66 \rightarrow 5392 \rightarrow 5391 \rightarrow 5388 \rightarrow 5291 \rightarrow 5278 \rightarrow 5289 \rightarrow 5290 \rightarrow 5283 \rightarrow 5284 \rightarrow 5280 \rightarrow T$

Shortest distance: 150335.55441905273

Total energy cost: 259087

Nodes expanded: 23284

Discussion:

For Task 2, we have used UCS again. We still have links of varying non-negative costs, but the conditions under which they are added to the priority queue have changed. Instead of maintaining a dictionary of path distances, we now maintain a dictionary of energy costs. A set of explored nodes is not maintained anymore as we may have to revisit a node since both the distance and the energy cost could change. Although a set of $\{node, distance, cost\}$ pairs could be maintained, we found it to be an overhead as it did not make any improvement to the number of nodes expanded. For each neighbour that is processed, we calculate the total distance and energy up till that point. We only add a neighbour to the queue and update its priority as long as the energy cost is within budget and the neighbour does not exist in the cost dictionary or the new cost is less than the existing one. Since the priority is still based on the distance, we will always expand the shortest nodes first. This ensures that we can get the shortest path within a budget, as we are first prioritising distance, and then optimizing the budget. Initially, we implemented a naive version of this solution, where the only difference from Task 1 was an extra check to ensure new nodes are added to the queue if and only if adding them does not exceed the energy budget. It gave us a path with a distance of 150784 and an energy cost of 287931. Although this distance is very close, it is not the optimal solution.

2.3 Task 3

Result: As expected, the shortest path, distance and total energy cost are the same as the UCS implementation in Task 2. However, the number of nodes expanded has reduced.

Nodes Expanded:

6991 using Manhattan distance with a weight of 0.76; 862 using Manhattan distance with a weight of 0.89 (Non-optimal alternative); 5270 (Weighted-Cost)

Weight Tuning:

Heuristic	Weight	Expanded	Reduction (%)	Weight	Expanded	Reduction (%)
Euclidean	0.15	21967	5.7	–	–	–
Manhattan	0.76	6991	70.0	0.89	862	96.3
Chebyshev	0.05	22949	1.4	–	–	–
Octile	0.11	22466	3.5	–	–	–

Discussion:

Task 3 requires improving the performance of Task 2 using a suitable heuristic. In terms of implementation, the only difference from Task 2 is that when adding a node to the priority queue, the priority is the sum of the distance and the heuristic cost ($g(n) + h(n)$). We have used Manhattan distance with a weight of 0.76 as it has the best performance amongst all the heuristics with 6991 expansions. No other heuristic comes even close as they have almost no performance gains to offer. The Manhattan heuristic works so well for this task that with a weight of 0.89, it can guide the naive version discussed earlier in Task 2 towards the correct solution in just 862 node expansions. This is a 96.3% reduction from the original 23284. This alternate solution does not guarantee optimality for all paths, but it works surprisingly well for the specific problem statement given to us. For Task 1, Manhattan was optimal at a weight of 0.76, but with the energy constraint, it is already optimal at 0.89. Perhaps the Manhattan heuristic is a good measure for cost and is biased towards it, which works out in our favour. We also implemented a version (*task3_weighted_energy*) where $h(n) = 0.1 * \text{energy_cost}$. It was able to find the correct path in 5270 node expansions, but more testing needs to be done to check whether this solution guarantees optimality.

3 Conclusions

In Task 1, we have seen how effective informed search can be when the right heuristic is used. The A* implementation with Octile distance as the heuristic expanded 88.5% fewer nodes than its uninformed UCS counterpart. Task 2 added an energy budget constraint to our shortest path problem. At first, it seemed like a simple additional check would do, but to implement it correctly, we had to flip our approach entirely. This also resulted in a much higher number of nodes being expanded, which is to say that even a simple constraint can make pathfinding more challenging. Yet again, a good heuristic can help improve performance, even when an additional constraint is at play. In Task 3, using Manhattan distance as the heuristic resulted in 87.7% less expansions than UCS. We also witnessed the classic trade-off between optimality and efficiency where a solution that does not guarantee optimality can achieve a solution that is close to optimal in much less time. Sometimes, getting a “good enough” solution quickly rather than the optimal solution slowly might be more desirable.